

截图 1

左上：监听通过 http 连接 sglang 的 ray 的 8000 端口

左下：监听直连 sglang 的 30000 端口

右：向 sglang 发起一次推理

```
tmzrz@localhost: ~  
[tmzrz@localhost:~]$ python3 /home/tmzrz/.cache/huggingface/hub/models--Qwen--Qwen3-0.6B/snapshots/c1899dc289a8d12100db378d81485cd75e07ca/v1/models/200  
[tmzrz@localhost:~]$ curl -s http://127.0.0.1:30000/v1/chat/completions -H "Content-Type: application/json" -d '{  
  "model": "Qwen/Qwen3-0.6B",  
  "messages": [{"role": "user", "content": "介绍一下你使用的模型及知识库的截止日期。"}],  
  "temperature": 0,  
  "chat_template_kwargs": {"enable_thinking": false}  
}' | python3 -c "import sys; json = json.loads(sys.stdin.read()); print(json.dumps(json.loads(sys.stdin.read())['choices'][0]['message']['content']))"  
tmzrz@localhost:~$ python3 /run_mooncake_workload.py  
Loaded 19 records from /home/tmzrz/conversation.filtered.jsonl  
Arrival rate: 2.0 req/s (Poisson)  
[a] status=200 in=6760 out=4 ttf=53.86ms latency=72.47ms  
[c] status=200 in=23141 out=0 ttf=Nonemes latency=83.3ms  
[7] status=200 in=26888 out=0 ttf=Nonemes latency=98.02ms  
[5] status=200 in=8834 out=167 ttf=497.41ms latency=1936.98ms  
[13] status=200 in=7324 out=15 ttf=157.12ms latency=595.09ms  
[2] status=200 in=7236 out=274 ttf=281.68ms latency=4658.86ms  
[10] status=200 in=11844 out=71 ttf=1358.37ms latency=3529.6ms  
[9] status=200 in=6758 out=338 ttf=689.11ms latency=7051.3ms  
[1] status=200 in=7322 out=339 ttf=511.2ms latency=6969.0ms  
[3] status=200 in=2290 out=380 ttf=42.93ms latency=6898.79ms  
[14] status=200 in=9418 out=145 ttf=102.93ms latency=3310.95ms  
[11] status=200 in=6324 out=258 ttf=1334.34ms latency=6109.02ms  
[8] status=200 in=10498 out=336 ttf=385.9ms latency=7068.93ms  
[9] status=200 in=17458 out=330 ttf=1635.20ms latency=7271.85ms  
[16] status=200 in=12846 out=305 ttf=98.35ms latency=4882.41ms  
[15] status=200 in=915 out=327 ttf=71.61ms latency=5361.83ms  
[12] status=200 in=2012 out=384 ttf=983.27ms latency=6737.37ms  
[17] status=200 in=20506 out=378 ttf=4230.41ms latency=8865.02ms  
[18] status=200 in=16609 out=350 ttf=8450.2ms latency=12269.78ms  
=====
```

截图 2

左：sglang 的运行状态

右：通过脚本完成的 Mooncake trace 采样 workload 的完成情况

```
tmzrz@localhost: ~  
[tmzrz@localhost:~]$ python3 /home/tmzrz/.cache/huggingface/hub/models--Qwen--Qwen3-0.6B/snapshots/c1899dc289a8d12100db378d81485cd75e07ca/v1/models/200  
[tmzrz@localhost:~]$ curl -s http://127.0.0.1:30000/v1/chat/completions -H "Content-Type: application/json" -d '{  
  "model": "Qwen/Qwen3-0.6B",  
  "messages": [{"role": "user", "content": "介绍一下你使用的模型及知识库的截止日期。"}],  
  "temperature": 0,  
  "chat_template_kwargs": {"enable_thinking": false}  
}' | python3 -c "import sys; json = json.loads(sys.stdin.read()); print(json.dumps(json.loads(sys.stdin.read())['choices'][0]['message']['content']))"  
tmzrz@localhost:~$ python3 /run_mooncake_workload.py  
Loaded 19 records from /home/tmzrz/conversation.filtered.jsonl  
Arrival rate: 2.0 req/s (Poisson)  
[a] status=200 in=6760 out=4 ttf=53.86ms latency=72.47ms  
[c] status=200 in=23141 out=0 ttf=Nonemes latency=83.3ms  
[7] status=200 in=26888 out=0 ttf=Nonemes latency=98.02ms  
[5] status=200 in=8834 out=167 ttf=497.41ms latency=1936.98ms  
[13] status=200 in=7324 out=15 ttf=157.12ms latency=595.09ms  
[2] status=200 in=7236 out=274 ttf=281.68ms latency=4658.86ms  
[10] status=200 in=11844 out=71 ttf=1358.37ms latency=3529.6ms  
[9] status=200 in=6758 out=338 ttf=689.11ms latency=7051.3ms  
[1] status=200 in=7322 out=339 ttf=511.2ms latency=6969.0ms  
[3] status=200 in=2290 out=380 ttf=42.93ms latency=6898.79ms  
[14] status=200 in=9418 out=145 ttf=102.93ms latency=3310.95ms  
[11] status=200 in=6324 out=258 ttf=1334.34ms latency=6109.02ms  
[8] status=200 in=10498 out=336 ttf=385.9ms latency=7068.93ms  
[9] status=200 in=17458 out=330 ttf=1635.20ms latency=7271.85ms  
[16] status=200 in=12846 out=305 ttf=98.35ms latency=4882.41ms  
[15] status=200 in=915 out=327 ttf=71.61ms latency=5361.83ms  
[12] status=200 in=2012 out=384 ttf=983.27ms latency=6737.37ms  
[17] status=200 in=20506 out=378 ttf=4230.41ms latency=8865.02ms  
[18] status=200 in=16609 out=350 ttf=8450.2ms latency=12269.78ms  
=====
```